



Representing Images



Why is Computer Vision so Hard?



Why is Computer Vision so Hard?

[illegible]

An image is just a grid, or **Matrix** of numbers between 0 and 1.

The number gives the **intensity**, or darkness, of a **Pixel**.

Why is Computer Vision so Hard?

0.8	0.7	0.7	0.7	0.7	0.8	0.7	0.7	0.8	0.9	0.6	0.6	0.9	1.0	1.0	0.9	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.8	0.9	0.7	0.6
0.7	0.7	0.8	0.8	0.8	0.7	0.7	0.6	0.7	0.9	1.0	0.5	0.2	0.6	0.9	1.0	0.9	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.9	0.8	0.6
0.7	0.8	0.8	0.7	0.6	0.4	0.4	0.3	0.4	0.8	1.0	1.0	0.6	0.1	0.2	0.5	0.8	0.9	0.8	0.7	0.7	0.8	0.7	0.7	0.7	0.7	0.8	0.7
0.7	0.7	0.7	0.4	0.2	0.1	0.1	0.2	0.5	0.9	1.0	1.0	0.7	0.2	0.1	0.1	0.3	0.7	0.8	0.7	0.7	0.7	0.7	0.8	0.8	0.7	0.8	0.7
0.7	0.6	0.3	0.1	0.1	0.0	0.1	0.4	0.8	1.0	1.0	1.0	0.7	0.2	0.1	0.1	0.1	0.3	0.6	0.7	0.7	0.7	0.7	0.8	0.8	0.7	0.8	0.8
0.8	0.2	0.1	0.1	0.1	0.3	0.6	0.9	1.0	1.0	0.9	0.6	0.3	0.1	0.1	0.1	0.1	0.1	0.5	0.7	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.9
0.3	0.1	0.1	0.1	0.1	0.2	0.5	0.7	0.9	1.0	1.0	0.9	0.5	0.3	0.1	0.1	0.0	0.3	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.7	0.7	0.8
0.1	0.1	0.1	0.1	0.2	0.3	0.5	0.7	0.9	1.0	0.9	0.8	0.4	0.3	0.1	0.0	0.0	0.0	0.1	0.6	0.8	0.7	0.7	0.7	0.8	0.8	0.9	0.8
0.2	0.2	0.2	0.2	0.4	0.4	0.6	0.7	0.9	0.9	0.8	0.5	0.4	0.3	0.2	0.1	0.0	0.0	0.3	0.7	0.8	0.8	0.8	0.8	0.8	0.8	0.9	1.0
0.3	0.3	0.3	0.3	0.4	0.6	0.7	0.9	0.8	0.6	0.4	0.4	0.4	0.3	0.2	0.1	0.0	0.0	0.1	0.7	0.8	0.8	0.8	0.8	0.7	0.8	0.9	1.0
0.3	0.4	0.3	0.3	0.4	0.6	0.7	0.9	0.7	0.5	0.4	0.3	0.3	0.3	0.2	0.1	0.0	0.0	0.0	0.5	0.8	0.8	0.8	0.8	0.7	0.7	0.8	0.9
0.4	0.4	0.4	0.4	0.4	0.6	0.6	0.7	0.6	0.5	0.4	0.3	0.3	0.3	0.2	0.2	0.1	0.0	0.0	0.5	0.7	0.7	0.7	0.7	0.7	0.7	0.8	0.9
0.4	0.5	0.4	0.4	0.4	0.5	0.6	0.6	0.6	0.5	0.4	0.4	0.3	0.3	0.3	0.2	0.1	0.0	0.1	0.4	0.7	0.7	0.7	0.7	0.7	0.7	0.7	0.8
0.5	0.5	0.5	0.4	0.5	0.5	0.7	0.6	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.2	0.1	0.2	0.4	0.7	0.7	0.7	0.7	0.6	0.6	0.6	0.6
0.6	0.6	0.6	0.5	0.5	0.7	0.6	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.7	0.6	0.3	0.1	0.2	0.3	0.6	0.6	0.6	0.6	0.6	0.6	0.6
0.8	0.8	0.7	0.6	0.5	0.5	0.6	0.6	0.5	0.4	0.4	0.4	0.5	0.4	0.6	0.9	0.8	0.4	0.1	0.2	0.3	0.6	0.6	0.5	0.5	0.5	0.5	0.6
0.9	0.9	0.8	0.6	0.5	0.5	0.5	0.4	0.4	0.4	0.4	0.4	0.5	0.5	0.8	0.8	0.8	0.5	0.2	0.2	0.3	0.5	0.5	0.5	0.5	0.5	0.5	0.5
1.0	1.0	1.0	0.9	0.8	0.5	0.5	0.4	0.4	0.4	0.4	0.3	0.4	0.6	0.9	0.9	0.7	0.5	0.2	0.1	0.3	0.4	0.4	0.4	0.4	0.4	0.4	0.5
1.0	1.0	0.9	0.8	0.6	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.6	0.9	0.9	0.7	0.2	0.0	0.3	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4
1.0	1.0	1.0	1.0	0.9	0.8	0.5	0.4	0.4	0.3	0.4	0.4	0.6	0.9	0.9	0.6	0.1	0.0	0.0	0.2	0.3	0.3	0.3	0.4	0.4	0.4	0.4	0.4
1.0	1.0	1.0	1.0	0.9	0.9	0.7	0.4	0.4	0.4	0.4	0.5	0.6	0.8	0.9	0.4	0.1	0.0	0.0	0.1	0.3	0.3	0.3	0.3	0.3	0.4	0.4	0.4
0.9	1.0	1.0	1.0	1.0	0.9	0.8	0.5	0.5	0.4	0.4	0.4	0.5	0.8	0.8	0.3	0.1	0.0	0.0	0.0	0.2	0.3	0.3	0.3	0.3	0.4	0.4	0.4
0.9	1.0	1.0	1.0	1.0	1.0	0.9	0.6	0.5	0.5	0.4	0.4	0.4	0.5	0.4	0.3	0.1	0.1	0.1	0.0	0.1	0.2	0.2	0.2	0.3	0.3	0.4	0.4
1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.7	0.6	0.5	0.6	0.5	0.5	0.3	0.3	0.4	0.8	0.6	0.5	0.6	0.6	0.5	0.4	0.2	0.2	0.3	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.7	0.6	0.6	0.7	0.7	0.5	0.4	0.5	0.7	0.6	0.3	0.2	0.3	0.4	0.5	0.7	0.4	0.3	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.8	0.8	0.8	0.5	0.4	0.6	0.5	0.7	0.8	0.8	0.8	0.8	0.8	0.4	0.3	0.3	0.3	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.5	0.3	0.4	0.7	0.8	0.9	0.9	0.9	0.9	0.9	0.8	0.8	0.5	0.4	0.3	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.5	0.3	0.3	0.8	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.4	0.6	0.3	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.5	0.2	0.3	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.4	0.5	0.3	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	1.0	0.9	0.4	0.1	0.3	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.6	0.5	0.3	
0.9	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.5	0.1	0.3	0.8	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.6	0.5	0.3	
0.9	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.6	0.4	0.5	0.9	0.8	0.9	0.9	0.9	0.9	0.8	0.7	0.8	0.5	0.3	
0.9	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.8	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.8	0.6	0.2	
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.8	0.7	0.2	

Can you see the Panda's head in this Matrix?

Why is Computer Vision so Hard?

0.8	0.7	0.7	0.7	0.7	0.8	0.7	0.7	0.8	0.9	0.6	0.6	0.9	1.0	1.0	0.9	0.7	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.8	0.9	0.7	0.6
0.7	0.7	0.8	0.8	0.8	0.7	0.7	0.6	0.7	0.9	1.0	0.5	0.2	0.6	0.9	1.0	0.9	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.8	0.9	0.8	0.6
0.7	0.8	0.8	0.7	0.6	0.4	0.4	0.3	0.4	0.8	1.0	1.0	0.6	0.1	0.2	0.5	0.8	0.9	0.8	0.7	0.7	0.8	0.7	0.7	0.7	0.7	0.8	0.8	0.7
0.7	0.7	0.7	0.4	0.2	0.1	0.1	0.2	0.5	0.9	1.0	1.0	0.7	0.2	0.1	0.1	0.3	0.7	0.8	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.7	0.8	0.7
0.7	0.6	0.3	0.1	0.1	0.0	0.1	0.4	0.6	1.0	1.0	1.0	1.0	0.7	0.2	0.1	0.1	0.3	0.6	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.7	0.8	0.9
0.8	0.2	0.1	0.1	0.1	0.3	0.6	0.9	1.0	1.0	0.9	0.6	0.3	0.1	0.1	0.1	0.1	0.5	0.7	0.7	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.9	0.8
0.3	0.1	0.1	0.1	0.1	0.2	0.5	0.7	0.9	1.0	1.0	0.9	0.5	0.3	0.1	0.1	0.0	0.3	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.8	0.7	0.8	0.9
0.1	0.1	0.1	0.1	0.2	0.3	0.5	0.7	0.9	1.0	0.9	0.8	0.4	0.3	0.1	0.0	0.0	0.0	1.0	0.6	0.8	0.7	0.7	0.8	0.8	0.8	0.7	0.8	0.9
0.2	0.2	0.2	0.2	0.4	0.4	0.6	0.7	0.9	0.9	0.8	0.5	0.4	0.3	0.2	0.1	0.0	0.0	0.3	0.7	0.8	0.8	0.8	0.8	0.8	0.8	0.9	1.0	1.0
0.3	0.3	0.3	0.3	0.4	0.6	0.7	0.9	0.8	0.6	0.4	0.4	0.4	0.3	0.2	0.1	0.1	0.0	0.1	0.7	0.8	0.8	0.8	0.8	0.7	0.8	0.9	1.0	1.0
0.3	0.4	0.4	0.3	0.3	0.4	0.6	0.7	0.9	0.7	0.5	0.4	0.3	0.3	0.2	0.1	0.0	0.0	0.0	0.5	0.8	0.8	0.8	0.8	0.7	0.7	0.8	0.9	0.9
0.4	0.4	0.4	0.4	0.4	0.6	0.6	0.7	0.6	0.5	0.4	0.3	0.3	0.3	0.2	0.2	0.1	0.0	0.0	0.5	0.7	0.7	0.7	0.7	0.7	0.7	0.8	0.8	0.9
0.4	0.5	0.4	0.4	0.4	0.5	0.6	0.6	0.6	0.5	0.4	0.4	0.3	0.3	0.3	0.2	0.1	0.0	0.1	0.4	0.7	0.7	0.7	0.7	0.7	0.7	0.7	0.7	0.8
0.5	0.5	0.5	0.4	0.5	0.5	0.7	0.6	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.2	0.1	0.2	0.4	0.7	0.7	0.7	0.7	0.7	0.6	0.6	0.6	0.6
0.6	0.6	0.6	0.5	0.5	0.7	0.6	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.7	0.6	0.3	0.1	0.2	0.3	0.6	0.6	0.6	0.6	0.6	0.6	0.6	0.6
0.8	0.8	0.7	0.6	0.5	0.5	0.6	0.6	0.5	0.4	0.4	0.4	0.5	0.4	0.6	0.9	0.8	0.4	0.1	0.2	0.3	0.6	0.6	0.5	0.5	0.5	0.5	0.6	0.6
0.9	0.9	0.8	0.6	0.5	0.5	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.5	0.5	0.8	0.8	0.5	0.2	0.2	0.3	0.5	0.5	0.5	0.5	0.5	0.5	0.5	0.5
1.0	1.0	1.0	0.9	0.8	0.6	0.5	0.3	0.4	0.4	0.4	0.3	0.4	0.6	0.9	0.9	0.7	0.5	0.2	0.1	0.3	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.5
1.0	1.0	1.0	0.9	0.8	0.6	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.6	0.9	0.9	0.7	0.2	0.0	0.3	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4
1.0	1.0	1.0	1.0	0.9	0.8	0.5	0.4	0.4	0.4	0.3	0.4	0.4	0.6	0.9	0.9	0.6	0.1	0.0	0.0	0.2	0.3	0.3	0.3	0.3	0.4	0.4	0.4	0.4
1.0	1.0	1.0	1.0	0.9	0.9	0.7	0.4	0.4	0.4	0.4	0.5	0.6	0.8	0.9	0.4	0.1	0.0	0.0	0.1	0.3	0.3	0.3	0.3	0.3	0.4	0.4	0.4	0.4
0.9	1.0	1.0	1.0	1.0	0.9	0.8	0.5	0.5	0.4	0.4	0.4	0.5	0.8	0.8	0.3	0.1	0.0	0.0	0.2	0.1	0.3	0.3	0.3	0.3	0.4	0.4	0.4	0.4
0.9	1.0	1.0	1.0	1.0	1.0	0.9	0.6	0.5	0.5	0.4	0.4	0.4	0.5	0.4	0.3	0.1	0.1	0.1	0.1	0.1	0.2	0.2	0.2	0.3	0.3	0.3	0.3	0.4
1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.7	0.6	0.5	0.6	0.5	0.5	0.3	0.3	0.3	0.4	0.6	0.5	0.6	0.6	0.5	0.4	0.2	0.2	0.3	0.3
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.7	0.6	0.6	0.7	0.7	0.5	0.4	0.5	0.7	0.6	0.3	0.2	0.3	0.4	0.5	0.7	0.4	0.3	0.3	0.3
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.8	0.8	0.8	0.5	0.4	0.6	0.5	0.7	0.8	0.8	0.8	0.8	0.8	0.4	0.3	0.3	0.3	0.3	0.3
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.5	0.3	0.4	0.7	0.8	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.6	0.5	0.3	0.3
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.5	0.3	0.3	0.8	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.4	0.5	0.3	0.3
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.9	0.9	0.9	0.5	0.2	0.3	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.9	0.8	0.7	0.3	0.2	0.2	0.2

Computer Vision is all about making sense of data like **this**.

Matrices and Arrays

Matrix is a generic term for a grid of numbers.

Array is a type of object in Python, which can represent Matrices.

Image is what we call it when we draw out arrays in a way that's easy to see.

We will use these terms mostly interchangeably.

Matrices and Arrays

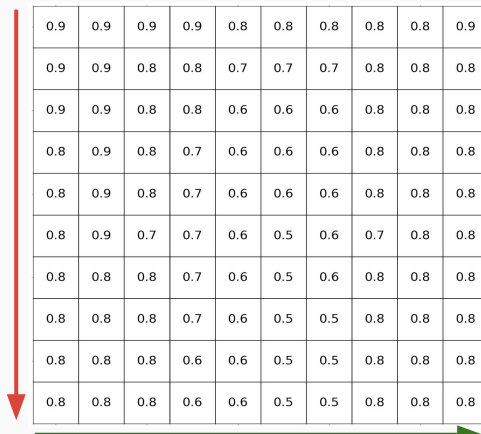
Arrays can have different numbers of **axes**, which describe how many rows and columns there are. For example, an array with **one axis** looks like this:

0.8	0.9	0.8	0.7	0.6	0.6	0.6	0.8	0.8	0.8
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

We call such an array a **vector**, or **1-dimensional array**.

Matrices and Arrays

An array with **two axes** looks like this:




A 10x10 grid representing a 2D array. The grid is composed of 10 rows and 10 columns of numerical values. To the left of the grid, a red arrow points downwards, labeled 'Rows'. Below the grid, a green arrow points to the right, labeled 'Columns'.

0.9	0.9	0.9	0.9	0.8	0.8	0.8	0.8	0.8	0.9
0.9	0.9	0.8	0.8	0.7	0.7	0.7	0.8	0.8	0.8
0.9	0.9	0.8	0.8	0.6	0.6	0.6	0.8	0.8	0.8
0.8	0.9	0.8	0.7	0.6	0.6	0.6	0.8	0.8	0.8
0.8	0.9	0.8	0.7	0.6	0.6	0.6	0.8	0.8	0.8
0.8	0.9	0.7	0.7	0.6	0.5	0.6	0.7	0.8	0.8
0.8	0.8	0.8	0.7	0.6	0.5	0.6	0.8	0.8	0.8
0.8	0.8	0.8	0.7	0.6	0.5	0.5	0.8	0.8	0.8
0.8	0.8	0.8	0.6	0.6	0.5	0.5	0.8	0.8	0.8
0.8	0.8	0.8	0.6	0.6	0.5	0.5	0.8	0.8	0.8

Matrices and Arrays

In fact, arrays can have even more than two axes! Here is an array **with 3 axes**:



Matrices and Arrays

We've already seen how a black and white image can be shown with a 2D array.

Can you think of something we might do with 3D arrays?


Representing Color



Pixels now vary in **color** and **saturation** as well as **intensity**.

Its no longer possible to express every pixel with just **ONE** number!


Representing Color

[illegible]

We can use 3D Arrays to store information about colors along the **third axis**.

RGB Color Model

Every pixel in an image is represented with 3 Numbers, telling us how
Red, Green, or Blue it is.




Colors can be purely red, green or blue.

Or merged to make different colors.


RGB Color Model

How can we express different shades of colors in RGB?




RGB Color Model

How can we express different shades of colors in RGB?









HSV Color Model



Hue, Saturation, Value is another popular format for colors.

This is nice because **all** the colors (hues) can be expressed on one axis.

		(0.0 , 1.0, 1.0)
		(0.3 , 1.0, 1.0)
		(0.6 , 1.0, 1.0)

We use the other two axes to tell how dark or saturated a pixel is.

Applications



0.9	0.9	0.9	0.9	0.8	0.8	0.8	0.8	0.8	0.9
0.9	0.9	0.8	0.8	0.7	0.7	0.7	0.8	0.8	0.8
0.9	0.9	0.8	0.8	0.6	0.6	0.6	0.8	0.8	0.8
0.8	0.9	0.8	0.7	0.6	0.6	0.6	0.8	0.8	0.8
0.8	0.9	0.8	0.7	0.6	0.6	0.6	0.8	0.8	0.8
0.8	0.9	0.7	0.7	0.6	0.5	0.6	0.7	0.8	0.8
0.8	0.8	0.8	0.7	0.6	0.5	0.6	0.8	0.8	0.8
0.8	0.8	0.8	0.7	0.6	0.5	0.5	0.8	0.8	0.8
0.8	0.8	0.8	0.6	0.6	0.5	0.5	0.8	0.8	0.8
0.8	0.8	0.8	0.6	0.6	0.5	0.5	0.8	0.8	0.8


Cropping an image involves **slicing** out rows and columns of the matrix you don't want.

Applications



Tinting an image requires “zeroing-out” channels for other colors.

Applications



Matrices can also be stacked, or **concatenated**, to create new images.

Recap

1. Every image is represented with a matrix, or an array with two or three axes.
2. Color Images use **3D arrays** to store information about colors along their third axis.
3. The **RGB** format describes every pixel with a **red**, **green**, and **blue** value.
4. Operations like **slicing** and **concatenation** let us manipulate images.